

COMP 4754

SQL

Fall 2024
Hafez Seliem

SQL Language

- By IBM in 1970
- Declarative!
 - Say **what** you want, not **how** to get it
- Two sublanguages:
 - DDL – Data Definition Language
 - Define and modify schema
 - DML – Data Manipulation Language
 - Queries can be written intuitively.
- RDBMS responsible for efficient evaluation.
 - Choose and run algorithms for declarative queries
 - Choice of algorithm must not affect query answer.

Example

- “BoatClub” database is to enable members of a boat club to reserve boats
 - Sailors —members of the boat club who reserve boats; and
 - Boats —boats in the club's inventory.

Sailors

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

Boats

bid	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

Reserves

sid	bid	day
1	102	9/12/2015
2	102	9/13/2015

The SQL DDL

```
CREATE TABLE sailors (  
  sid    INTEGER,  
  sname  CHAR(20),  
  rating INTEGER,  
  age    REAL,  
  PRIMARY KEY (sid));
```


```
CREATE TABLE Boats (  
  bid    INTEGER,  
  bname  CHAR(20),  
  color  CHAR(10),  
  PRIMARY KEY (bid));
```

```
CREATE TABLE Reserves (  
  sid INTEGER,  
  bid INTEGER,  
  day DATE,  
  PRIMARY KEY (sid, bid, day),  
  FOREIGN KEY (sid) REFERENCES  
  sailors(sid),  
  FOREIGN KEY (bid) REFERENCES  
  Boats(bid));
```

<u>sid</u>	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

<u>bid</u>	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13



MySQL data type categories

- Character
- Numeric
- Date and time
- Binary
- Large Object (LOB)
- Spatial
- JSON

The character types

Type	Bytes
CHAR(M)	Mx4
VARCHAR(M)	L+1

- How the character types work (in case of default character set utf8mb4) ?

Data type	Original value	Value stored Bytes
CHAR (2)	'CA'	8
CHAR (10)	'CA'	40
VARCHAR (10)	'CA'	3
VARCHAR (20)	'California'	11
VARCHAR (20)	'New York' 'New York'	9

The integer types

Type	Bytes
BIGINT	8
INT	4
MEDIUMINT	3
SMALLINT	2
TINYINT	1

The fixed-point type

- **DECIMAL(M, D)**
 - M total number of digits
 - D number of digits in the fractional part.
- **The floating-point types**

Type	Bytes
DOUBLE	8
FLOAT	4

- **How the fixed-point and floating-point types work ?**

Data type	Original value	Value stored	Bytes used
DECIMAL(9,2)	1.20	1.20	5
DECIMAL(9,2)	1234567.89	1234567.89	5
DECIMAL(9,2)	-1234567.89	-1234567.89	5
DECIMAL(18,9)	1234567.89	1234567.89	8
DOUBLE	1234567.89	1234567.89	8
FLOAT	1234567.89	1234570.00	4

The date and time types

Type	Bytes
DATE	3
TIME	3
DATETIME	8
TIMESTAMP	4
YEAR	1

- The **DATETIME** column does absolutely nothing concerning time zones. **TIMESTAMP** always converting to and from UTC.

```
CREATE TABLE timezone_test ( `timestamp` TIMESTAMP,  
`datetime` DATETIME );
```

Example

```
SET SESSION time_zone = '+00:00';  
INSERT INTO timezone_test VALUES ('2029-02-14 08:47', '2029-02-14 08:47');  
SELECT * FROM timezone_test;
```

Output:

```
-- | timestamp | datetime |  
-- |-----|-----|  
-- | 2029-02-14 08:47:00 | 2029-02-14 08:47:00 |
```

```
SET SESSION time_zone = '-05:00';  
SELECT * FROM timezone_test;
```

Output:

```
-- | timestamp      | datetime      |  
-- |-----|-----|  
-- | 2029-02-14 03:47:00 | 2029-02-14 08:47:00 |
```

How MySQL interprets literal date/time values

Literal value	Stored Value
'2022-08-15'	2022-08-15
'2022-8-15'	2022-08-15
'22-8-15'	2022-08-15
'20220815'	2022-08-15
20220815	2022-08-15
'8-15-22'	(error)
'2022-02-31'	(error)

How MySQL interprets literal date/time values

-

Literal value	Stored Value
'7:32'	7:32:00
'19:32:11'	19:32:11
'193211'	19:32:11
193211	19:32:11
'19:61:11'	(error)

-

DATETIME or TIMESTAMP	Stored Value
'2022-08-15 19:32:11'	2022-08-15 19:32:11
'2022-08-15'	2022-08-15 00:00:00

How an ENUM('Yes','No','Maybe') column works

Value	Value stored	Value displayed
'Yes'	1	'Yes'
'No'	2	'No'
'Maybe'	3	'Maybe'
'Possibly'	(error)	
"	(error)	

How a SET ('Pepporoni','Mushrooms','Olives') column works ?

•Value	Valued stored (binary)	Value displayed
'Pepperoni'	1 (00000001)	'Pepperoni'
'Mushrooms'	2 (00000010)	'Mushrooms'
'Olives'	4 (00000100)	'Olives'
'Olives,Pepperoni'	5 (00000101)	'Pepperoni,Olives'
'Olives,Olives,Mushrooms'	6 (00000110)	'Mushrooms,Olives'
' Pepperoni , Sausage '	(error)	

The binary data types

Type	Bytes
BINARY(M)	M
VARCHAR(M)	L+1

The large object types

Type	Bytes
LONGBLOB	L+4
MEDIUMBLOB	L+3
BLOB	L+2
TINYBLOB	L+1
LONGTEXT	L+4
MEDIUMTEXT	L+3
TEXT	L+2
TINYTEXT	L+1

- **Terms to know about large objects**
 - BLOB (binary large object) types
 - CLOB (character large object) types

Constraints

- Recall that the schema defines the legal instances of the relations.
- Data types are a way to limit the kind of data that can be stored in a table, but they are often insufficient.
 - e.g., prices must be positive values
 - uniqueness, referential integrity, etc.
- Can specify constraints on individual columns or on tables.

Integrity Constraints

- IC conditions that every legal instance of a relation must satisfy.
 - Inserts/deletes/updates that violate ICs are disallowed.
 - Can ensure application semantics (e.g., sid is a key),
...or prevent inconsistencies (e.g., sname has to be a string,
age must be < 200)
- **Types of IC's:** Domain constraints, primary key constraints, foreign key constraints, general constraints.
 - **Domain constraints:** Field values must be of right type.
Always enforced.

Primary Keys

- A set of fields is a **superkey** if:
 - No two distinct tuples can have same values in all these fields
- A set of fields is a **key** for a relation if it is minimal:
 - It is a superkey
 - No subset of the fields is a superkey
- what if >1 key for a relation?
 - One of the keys is chosen to be the **primary key**. Other keys are called **candidate keys**.
- For example:
 - sid is a key for Students.
 - What about name?
 - The set {sid, gpa} is a superkey.

Primary and Candidate Keys

- Possibly many candidate keys (specified using UNIQUE), one of which is chosen as the primary key.
 - Keys must be used carefully!

```
CREATE TABLE Enrolled1
(sid    CHAR(20),
cid    CHAR(20),
grade  CHAR(2),
PRIMARY KEY
(sid,cid))
```

Not good either!

```
CREATE TABLE Enrolled2
(sid    CHAR(20),
cid    CHAR(20),
grade  CHAR(2),
PRIMARY KEY (sid),
UNIQUE (cid,
grade))
```

“For a given student and course, there is a single grade.”

Foreign Keys, Referential Integrity

- Foreign key: a “logical pointer”
 - Set of fields in a tuple in one relation that ‘refer’ to a tuple in another relation.
 - Reference to *primary* key of the other relation.

Foreign Keys in SQL

- For example, only students listed in the Students relation should be allowed to enroll for courses.
 - sid is a foreign key referring to Students:


```
CREATE TABLE Enrolled  
(sid CHAR(20), cid CHAR(20), grade CHAR(2),  
PRIMARY KEY (sid,cid),  
FOREIGN KEY (sid) REFERENCES Students(sid));
```

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B
11111	English102	A

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8



Enforcing Referential Integrity

- *sid* in Enrolled: foreign key referencing Students.
- **Scenarios:**
 - Insert Enrolled tuple with non-existent student id?
 - Delete a Students tuple?
 - Also delete Enrolled tuples that refer to it? (CASCADE)
 - Disallow if referred to? (NO ACTION)
 - Set *sid* in referring Enrolled tups to a default value? (SET DEFAULT)
 - Set *sid* in referring Enrolled tuples to null, denoting 'unknown' or 'inapplicable'. (SET NULL)
- Similar issues arise if primary key of Students tuple is updated.

Foreign keys actions

```
CREATE TABLE Enrolled  
(sid CHAR(20), cid CHAR(20), grade CHAR(2),  
PRIMARY KEY (sid,cid),  
FOREIGN KEY (sid) REFERENCES Students(sid)  
ON DELETE NO ACTION );
```

VS

```
FOREIGN KEY (sid) REFERENCES Students(sid)  
ON DELETE CASCADE);
```

VS

```
FOREIGN KEY (sid) REFERENCES Students(sid)  
ON DELETE SET NULL);
```


SQL: Modification Commands

- **Deletion:** `DELETE FROM <relation>`
`[WHERE <predicate>]`

- Example:

1. `DELETE FROM account where`
`balance > 100`
`-- deletes all tuples with balance > 100`

2. `DELETE FROM account`
`WHERE bname IN (SELECT bname`
`FROM branch`
`WHERE bcity = 'Bkln')`
`-- deletes all accounts from Brooklyn branch`

`account(bname, acct_no, balance)`

SQL: Modification Commands

Insertion:

INSERT INTO <relation> values (... , ... , ...)

or

INSERT INTO <relation>(att1, ..., attn) values(..., ..., ...)

or

INSERT INTO <relation> <query expression>

SQL: Modification Commands

Insertion:

Examples:

```
INSERT INTO account VALUES ( 'Perry' , A-768, 1200)
```

or

```
INSERT INTO account( bname, acct_no, balance)  
VALUES ( 'Perry' , A-768, 1200)
```

```
INSERT INTO account  
SELECT  bname, lno, 200  
FROM    loan  
WHERE   bname = 'Kenmore'
```

gives free \$200 savings account for each loan holder at Kenmore

SQL: Modification Commands

Update:

```
UPDATE <relation>  
    SET <attribute> = <expression>  
    WHERE <predicate>
```

SQL: Modification Commands

- Example:

```
UPDATE account
  SET    balance = balance * 1.06
 WHERE  balance > 10000
```

```
UPDATE account
  SET    balance = balance * 1.05
 WHERE  balance <= 10000
```

Alternative:

```
UPDATE account
  SET    balance =
          (CASE
            WHEN balance <= 10000 THEN balance*1.05
            ELSE balance*1.06
          END)
```

SQL DML 1: Basic Single-Table Queries

- **SELECT** [**DISTINCT**] *<column expression list>*
 FROM *<single table>*
 [**WHERE** *<predicate>*]
 [**GROUP BY** *<column list>*
 [**HAVING** *<predicate>*]
 [**ORDER BY** *<column list>*] ;

Basic Single-Table Queries

- **SELECT** [**DISTINCT**] *<column expression list>*
FROM *<single table>*
[**WHERE** *<predicate>*]
[**GROUP BY** *<column list>*
[**HAVING** *<predicate>*]
[**ORDER BY** *<column list>*] ;
- Simplest version is straightforward
 - Produce all tuples in the table that satisfy the predicate
 - Output the expressions in the SELECT list
 - Expression can be a column reference, or an arithmetic expression over column refs

Basic Single-Table Queries

- Example
 - **SELECT** S.name, S.gpa
FROM students S
WHERE S.dept = 'CS'
- Simplest version is straightforward
 - Produce all tuples in the table that satisfy the predicate
 - Output the expressions in the SELECT list
 - Expression can be a column reference, or an arithmetic expression over column refs

Basic Single-Table Queries

- Example:
 - `SELECT DISTINCT S.name, S.gpa`
`FROM students S`
`WHERE S.dept = 'CS';`
- DISTINCT flag specifies removal of duplicates before output

ORDER BY

- `SELECT DISTINCT S.name, S.gpa
FROM students S
WHERE S.dept = 'CS'
ORDER BY S.gpa DESC, S.name ASC;`
- Ascending order by default, but can be overridden
 - DESC flag for descending, ASC for ascending

Aggregates

- `SELECT AVG(S.gpa)`
 `FROM students S`
 `WHERE S.dept = 'CS'`
- Before producing output, compute a summary (a.k.a. an aggregate) of some arithmetic expression
- Produces 1 row of output
 - with one column in this case
- Other aggregates: SUM, COUNT, MAX, MIN
- Note: can use DISTINCT inside the agg function
 - `SELECT COUNT(DISTINCT S.name) FROM Students S`
 - vs. `SELECT DISTINCT COUNT (S.name) FROM Students S;`